

```
In [1]: # Import required libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import nltk
import re
import string
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer

# NLTK downloads (only needed once)
nltk.download('stopwords')
nltk.download('wordnet')

# Load the dataset
file_path = r"C:\Users\AkshS\OneDrive\Desktop\Imdb.xlsx"
df = pd.read_excel(file_path)

# Display the first few rows
print("Sample Data:")
print(df.head())

# Basic info
print("\nDataset Info:")
print(df.info())

# Check for missing values
print("\nMissing Values:")
print(df.isnull().sum())

# Class distribution
print("\nSentiment Distribution:")
print(df['sentiment'].value_counts())

# Review length analysis
df['review_length'] = df['review'].apply(lambda x: len(str(x).split()))
print("\nReview Length Description:")
print(df['review_length'].describe())

# Visualization: Sentiment count
sns.countplot(x='sentiment', data=df, palette='Set2')
plt.title('Sentiment Class Distribution')
plt.show()

# Visualization: Review Lengths
sns.histplot(df['review_length'], bins=30, kde=True)
plt.title('Review Length Distribution')
plt.xlabel('Number of Words')
plt.ylabel('Frequency')
plt.show()

# Clean the reviews
stop_words = set(stopwords.words('english'))
lemmatizer = WordNetLemmatizer()

def clean_text(text):
    text = text.lower()
    text = re.sub(r'<.*?>', '', text) # remove HTML
```

```

text = text.translate(str.maketrans('', '', string.punctuation)) # remove p
text = re.sub(r'\d+', '', text) # remove numbers
words = text.split()
words = [lemmatizer.lemmatize(word) for word in words if word not in stop_wc
return ' '.join(words)

df['clean_review'] = df['review'].astype(str).apply(clean_text)

print("\nOriginal vs Cleaned:")
print(df[['review', 'clean_review']].head())

```

```

[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\AkshS\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\AkshS\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!

```

Sample Data:

	review	sentiment
0	One of the other reviewers has mentioned that ...	positive
1	A wonderful little production. The...	positive
2	I thought this was a wonderful way to spend ti...	positive
3	Basically there's a family where a little boy ...	negative
4	Petter Mattei's "Love in the Time of Money" is...	positive

Dataset Info:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   review      50000 non-null  object
1   sentiment   50000 non-null  object
dtypes: object(2)
memory usage: 781.4+ KB
None

```

Missing Values:

```

review      0
sentiment   0
dtype: int64

```

Sentiment Distribution:

```

sentiment
positive    25000
negative    25000
Name: count, dtype: int64

```


Review Length Description:


```

count      50000.000000
mean        231.137900
std         171.339334
min           1.000000
25%         126.000000
50%         173.000000
75%         280.000000
max         2470.000000
Name: review_length, dtype: float64

```

```
C:\Users\AkshS\AppData\Local\Temp\ipykernel_9176\516829515.py:41: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v  
0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.  
  
sns.countplot(x='sentiment', data=df, palette='Set2')
```





Original vs Cleaned:

	review \
0	One of the other reviewers has mentioned that ...
1	A wonderful little production. The...
2	I thought this was a wonderful way to spend ti...
3	Basically there's a family where a little boy ...
4	Petter Mattei's "Love in the Time of Money" is...

	clean_review
0	one reviewer mentioned watching oz episode you...
1	wonderful little production filming technique ...
2	thought wonderful way spend time hot summer we...
3	basically there family little boy jake think t...
4	petter matteis love time money visually stunni...

```
In [2]: from sklearn.feature_extraction.text import TfidfVectorizer

# TF-IDF vectorization
vectorizer = TfidfVectorizer(max_features=5000)
X = vectorizer.fit_transform(df['clean_review'])

print("TF-IDF Matrix Shape:", X.shape)

# Add simple textual features
df['char_count'] = df['review'].apply(lambda x: len(str(x)))
df['avg_word_length'] = df['char_count'] / df['review_length']

print("\nTextual Feature Samples:")
print(df[['review_length', 'char_count', 'avg_word_length']].head())
```

TF-IDF Matrix Shape: (50000, 5000)

Textual Feature Samples:

	review_length	char_count	avg_word_length
0	307	1761	5.736156
1	162	998	6.160494
2	166	926	5.578313
3	138	748	5.420290
4	230	1317	5.726087

```
In [5]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import LinearSVC

# Convert target labels
y = df['sentiment']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

# Logistic Regression
lr = LogisticRegression()
lr.fit(X_train, y_train)

# Naive Bayes
nb = MultinomialNB()
nb.fit(X_train, y_train)

# Support Vector Machine
svm = LinearSVC()
svm.fit(X_train, y_train)
```

```
Out[5]: LinearSVC
LinearSVC()
```

```
In [7]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import seaborn as sns

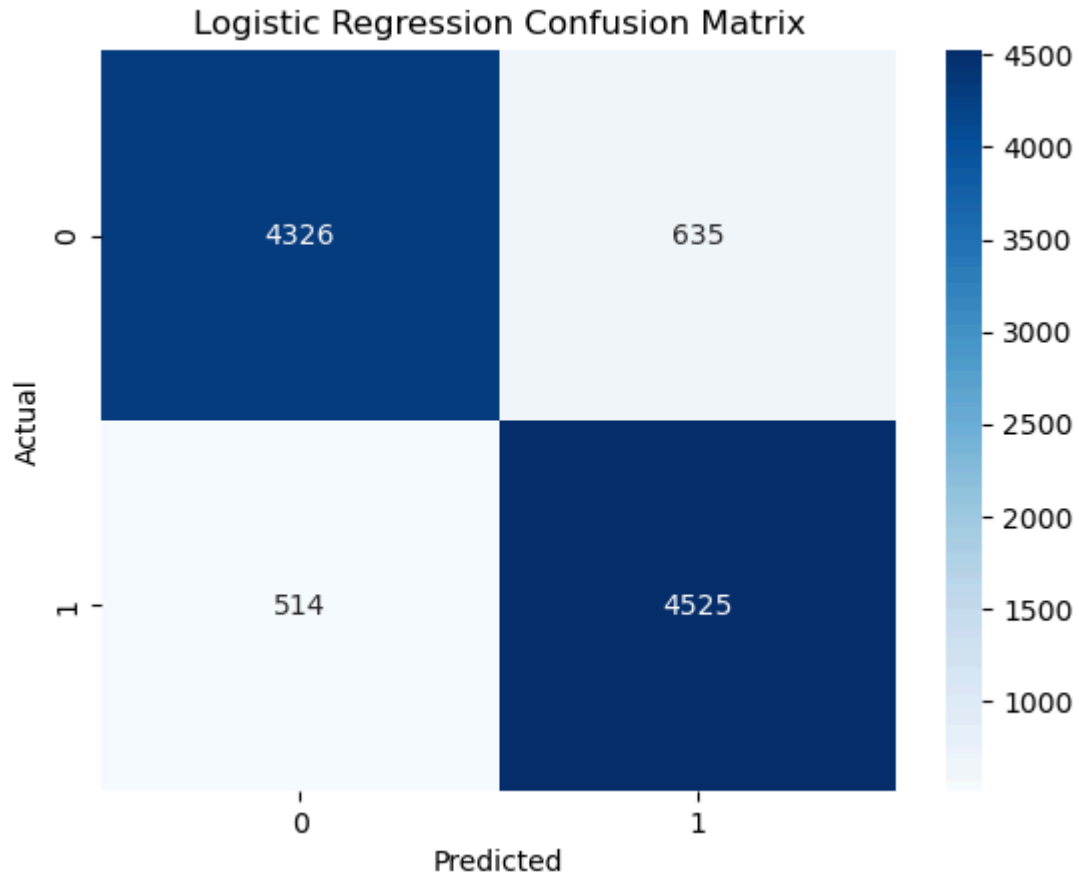
# Function to evaluate models
def evaluate_model(model, name):
    print(f"\n{name} Model Evaluation:")
    y_pred = model.predict(X_test)
    print(classification_report(y_test, y_pred))
    print("Accuracy:", accuracy_score(y_test, y_pred))
    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title(f'{name} Confusion Matrix')
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.show()

# Evaluate all models
evaluate_model(lr, "Logistic Regression")
evaluate_model(nb, "Naive Bayes")
evaluate_model(svm, "SVM")
```

Logistic Regression Model Evaluation:

	precision	recall	f1-score	support
negative	0.89	0.87	0.88	4961
positive	0.88	0.90	0.89	5039
accuracy			0.89	10000
macro avg	0.89	0.88	0.89	10000
weighted avg	0.89	0.89	0.89	10000


Accuracy: 0.8851



Naive Bayes Model Evaluation:

	precision	recall	f1-score	support
negative	0.85	0.84	0.85	4961
positive	0.85	0.85	0.85	5039
accuracy			0.85	10000
macro avg	0.85	0.85	0.85	10000
weighted avg	0.85	0.85	0.85	10000


Accuracy: 0.8487



SVM Model Evaluation:

	precision	recall	f1-score	support
negative	0.88	0.87	0.88	4961
positive	0.87	0.89	0.88	5039
accuracy			0.88	10000
macro avg	0.88	0.88	0.88	10000
weighted avg	0.88	0.88	0.88	10000

Accuracy: 0.8786



In []:

In []: